

NEEL SHIRISH MORE

MACHINE LEARNING ENGINEER

United Kingdom | +44 7393 063 988 | neelmore007@gmail.com

[Linkedin](#) | [Github](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Experience in developing and optimising machine learning models using Python and frameworks such as PyTorch, with strong focus on performance, scalability, and model evaluation. Works on large-scale, distributed training and inference systems, identifying bottlenecks and improving efficiency across compute-intensive workloads. Designs and conducts experiments on advanced AI methods, contributing to model optimisation and benchmarking. Collaborates with research and software teams to build and test production-grade AI systems, with interest in efficient compute, model scaling, and next-generation AI infrastructure. Applies strong software engineering practices, with ability to communicate complex technical concepts and contribute to high-performance AI systems.

TECHNICAL SKILLS

Deep Learning & Frameworks: PyTorch, Hugging Face Transformers, XGBoost, scikit-learn, CNNs, Diffusion Models, LLMs.

Model Training & Optimisation: Distributed Training Concepts, Low-Precision Arithmetic, Model Profiling, Performance Benchmarking, Hyperparameter Tuning, Evaluation.

Languages: Python, C++, SQL, Bash.

MLOps & Infrastructure: Docker, CI/CD, Experiment Tracking (MLFlow), Model Serialisation, Linux, Cloud Deployment.

Engineering Practices: Code Review, Experiment Design, Cross-Functional Collaboration, Technical Communication, Agile.

EXPERIENCE

Machine Learning Engineer

RTN Zero Consulting

Dec 2025 - Present

Manchester, UK

- Designed and trained a 6-stage chained ML pipeline (XGBoost, PyTorch) for multi-tenant forecasting — optimising models for throughput and accuracy across sales, inventory, operations, and carbon emissions domains.
- Benchmarked ML inference runtimes (MLX vs Ollama) under controlled conditions — measuring token throughput, memory efficiency, and thermal stability to inform model selection for production deployment.
- Built multi-agent LLM systems with tool-calling orchestration, implementing grounding and response validation to improve reliability — evaluated performance across different model architectures and configurations.
- Developed automated model retraining pipelines with data quality guardrails and evaluation checks — ensuring consistent performance across deployment scenarios and data conditions.

Data Analyst

Subtle Solutions

Sept 2023 - Oct 2025

India (Remote)

- Built scalable data pipelines, visualisation dashboards, and ETL workflows to handle high-volume operational data — delivering data-driven insights through automated reporting and API integrations.
- Collaborated with cross-functional teams to translate business requirements into technical solutions — applying iterative delivery practices and maintaining code quality through review and testing.

Software Engineering Intern

JPMorgan Chase & Co.

Dec 2021

Remote

- Developed data processing pipelines for financial data — implementing validation, transformation, and quality checks for reliable downstream analytics.

PROJECTS

Adaptive Inference in Masked Diffusion Models | *PyTorch, Python, Research Implementation*

- Reproduced and extended masked diffusion model research (ICML 2025) — implementing adaptive token ordering to investigate how inference-time strategies affect model reliability and efficiency in discrete domains.
- Designed and ran experiments comparing sampling strategies across reasoning benchmarks (Sudoku, logical puzzles), analysing performance characteristics and compute trade-offs.

Fine-Tune GPT-2: LLM Training & Optimisation Pipeline | *PyTorch, Hugging Face Transformers, MLFlow, Docker*

- Built end-to-end LLM fine-tuning pipelines with experiment tracking and hyperparameter optimisation — managing data preparation, model training, evaluation, and reproducible workflows.
- Conducted systematic evaluations of fine-tuned model behaviour under different training configurations — analysing data composition effects on output quality and model reliability.

Multi-Document Summarisation System | *Hugging Face Transformers, FastAPI, Docker, scikit-learn*

- Built a hybrid multi-document summarisation pipeline using Sentence-Transformers for semantic embeddings and KMeans clustering — deployed as a FastAPI service with Docker for reproducible inference.
- Evaluated summarisation quality across transformer architectures (BART, Pegasus), benchmarking latency and output coherence against baseline methods.

EDUCATION

MSc. Artificial Intelligence

University of Manchester

Sept 2023 – Sept 2024

Manchester, UK

- **Core Focus:** Machine Learning, Deep Learning, Large-Scale Systems, Computer Vision, Data Science
- **Completed hands-on modules in model training, evaluation, and optimisation — building and benchmarking ML models across supervised, unsupervised, and deep learning paradigms.**

B.E. Computer Engineering

Bharati Vidyapeeth's College of Engineering

Aug 2019 – Aug 2023

Pune, India

- **Core Focus:** Data Engineering, Operating Systems, Algorithms, NLP, Network Security
- **First-class with distinction (81%).**

PUBLICATIONS

- Kadam Pratima, More Neel. et al. (2023) 'Deep Learning and Artificial Intelligence for Games', IOSR Journal of Computer Engineering.